


EML'24 – Lecture 5

Classification II

ISLR 4, ESL 4

Prof. Isabel Valera
14 November 2024



Fitting Classification Models



Fitting Logistic Regression Models

Recall that the multivariate logistic regression model is defined as

- $\log\left(\frac{p(Y=1|X)}{1-p(Y=1|X)}\right) = \beta_0 + \beta_1 X + \dots + \beta_p X_p$ with $p(Y = 1|X) = \frac{e^{\beta_0 + \beta_1 X + \dots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X + \dots + \beta_p X_p}}$

We usually fit a logistic regression model by maximum likelihood

- log-likelihood function $\ell(\theta) = \sum_{i=1}^n \log p_{g_i}(x_i; \theta)$ and density function $p_k(x_i, \theta) = \Pr(G = k | X = x_i; \theta)$
- for a binary problem, class coding $y_i = \begin{cases} 1 & | \ g_i = 1 \\ 0 & | \ g_i = 0 \end{cases}$ gives us $p_1(x; \theta) = p(x; \theta)$ and $p_0(x; \theta) = 1 - p(x; \theta)$

The log-likelihood then becomes

$$\ell(\boldsymbol{\beta}) = \sum_{i=1}^n \{y_i \log p(x_i; \boldsymbol{\beta}) + (1 - y_i) \log(1 - p(x_i; \boldsymbol{\beta}))\} = \sum_{i=1}^n \{y_i \boldsymbol{\beta}^T x_i - \log(1 + e^{\boldsymbol{\beta}^T x_i})\}$$

- where $\boldsymbol{\beta} = \{\beta_0, \beta_1, \dots\}$ and x_i a vector of the input values padded with a constant term $X_0 = 1$



Side calculation

$$\ell(\boldsymbol{\beta}) = \sum_{i=1}^n \{y_i \log p(x_i; \boldsymbol{\beta}) + (1 - y_i) \log(1 - p(x_i; \boldsymbol{\beta}))\}$$

$$= \sum_{i=1}^n \left\{ y_i \log \frac{e^{\beta_0 + \beta_1^T x_i}}{1 + e^{\beta_0 + \beta_1^T x_i}} + (1 - y_i) \log \frac{1}{1 + e^{\beta_0 + \beta_1^T x_i}} \right\} \quad (\text{definition of } p(x_i; \boldsymbol{\beta}))$$

$$= \sum_{i=1}^n \left\{ y_i \left[(\beta_0 + \beta_1^T x_i) - \log(1 + e^{\beta_0 + \beta_1^T x_i}) \right] - (1 - y_i) \log(1 + e^{\beta_0 + \beta_1^T x_i}) \right\} \quad (\log a/b = \log a - \log b)$$

$$= \sum_{i=1}^n \left\{ y_i \boldsymbol{\beta}^T x_i - \log(1 + e^{\boldsymbol{\beta}^T x_i}) \right\} \quad (\text{simplify})$$



Fitting Logistic Regression Models

We find the $\boldsymbol{\beta}$ that achieves maximum likelihood by setting the derivative to zero

- this yields the **score equations**

$$\frac{\partial \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = \sum_{i=1}^n x_i (y_i - p(x_i; \boldsymbol{\beta})) = 0$$

- these can be broken down to $p + 1$ equations that are **nonlinear** in $\boldsymbol{\beta}$
- because the first value of x_i is 1, the first equation takes the shape

$$\sum_{i=1}^n y_i = \sum_{i=1}^n p(x_i; \boldsymbol{\beta})$$

- the expected number of class-1 assignments is the number class-1 we observed



Fitting Logistic Regression Models

We can solve the score equations numerically using Newton-Raphson

$$\boldsymbol{\beta}^{new} = \boldsymbol{\beta}^{old} - \left(\frac{\partial^2 \ell(\boldsymbol{\beta}^{old})}{\partial \boldsymbol{\beta} \partial \boldsymbol{\beta}^T} \right)^{-1} \frac{\partial \ell(\boldsymbol{\beta}^{old})}{\partial \boldsymbol{\beta}}$$

- i.e. adjust coefficients proportionally to **second derivative** in the **opposite** direction of **first derivative**
- repeat until convergence
- note that $\frac{\partial^2 \ell(\boldsymbol{\beta}^{old})}{\partial \boldsymbol{\beta} \partial \boldsymbol{\beta}^T} = -\sum_{i=1}^n x_i x_i^T p(x_i; \boldsymbol{\beta})(1 - p(x_i; \boldsymbol{\beta}))$ is our old friend, the Hessian matrix!

Log-likelihood is **concave**

- single starting point suffices, $\boldsymbol{\beta} = \mathbf{0}$ is fine
- typically converges, but overshooting can occur
- diagonal of the Hessian matrix contains the squared standard deviations of outputs in the training set



Fitting Logistic Regression Models

In matrix notation we have

$$\frac{\partial \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = \mathbf{X}^T (\mathbf{y} - \mathbf{p}) \quad \frac{\partial^2 \ell(\boldsymbol{\beta})}{\partial \boldsymbol{\beta} \partial \boldsymbol{\beta}^T} = -\mathbf{X}^T \mathbf{W} \mathbf{X}$$

- where \mathbf{W} is a diagonal matrix with elements $w_{ii} = -p(x_i; \boldsymbol{\beta}^{old}) (1 - p(x_i; \boldsymbol{\beta}^{old}))$

A single Newton-Raphson step is

$$\begin{aligned} \boldsymbol{\beta}^{new} &= \boldsymbol{\beta}^{old} - (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{p}) = (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} (\mathbf{X} \boldsymbol{\beta}^{old} - \mathbf{W}^{-1} (\mathbf{y} - \mathbf{p})) = (\mathbf{X}^T \mathbf{W} \mathbf{X})^{-1} \mathbf{X}^T \mathbf{W} \mathbf{z} \\ \mathbf{z} &= \mathbf{X} \boldsymbol{\beta}^{old} - \mathbf{W}^{-1} (\mathbf{y} - \mathbf{p}) \end{aligned}$$

- a linear least-squares problem with output \mathbf{z} weighted by diagonal matrix \mathbf{W}

$$\boldsymbol{\beta}^{new} = \arg \min_{\boldsymbol{\beta}} (\mathbf{z} - \mathbf{X} \boldsymbol{\beta})^T \mathbf{W} (\mathbf{z} - \mathbf{X} \boldsymbol{\beta})$$

Linear Discriminant Analysis

The Bayes-optimal choice is to classify \mathbf{x} to the class with the largest discriminant

- the **discriminant** of a class k is the log-probability that cancels in the log-odds

$$\log\left(\frac{p_k(\mathbf{x})}{p_l(\mathbf{x})}\right) = \delta_k(\mathbf{x}) - \delta_l(\mathbf{x})$$

- Where we assume the class-conditional densities to be Gaussian, yielding:

$$\delta_k(\mathbf{x}) = \mathbf{x} \cdot \frac{\boldsymbol{\mu}_k}{\sigma^2} - \frac{\boldsymbol{\mu}_k^2}{2\sigma^2} + \log \pi_k$$

is the log-numerator from previous slide with the class-independent terms removed

Fitting Univariate LDA Models

In general, we do not know the underlying class densities but assume they are Gaussians, so that

- we estimate these using the finite training sample

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i: y_i=k} x_i$$


$$\hat{\sigma}^2 = \frac{1}{n-K} \sum_{k=1}^K \sum_{i: y_i=k} (x_i - \hat{\mu}_k)^2$$

$$\pi = n_k/n$$

- we assign x to the class with the largest fitted discriminant

$$\hat{\delta}_k(x) = x \cdot \frac{\hat{\mu}_k}{\hat{\sigma}^2} - \frac{\hat{\mu}_k^2}{2\hat{\sigma}^2} + \log \hat{\pi}_k$$

- **note** that the discriminants are linear (!)



*LDA fit over 20 samples per class,
fitted decision boundary in dashed black.
Bayes error 10.6%, LDA test error 11.1%*



Fitting LDA and QDA Models

Again, we use sample estimates

- $\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k} x_i$
- $\hat{\Sigma} = \frac{1}{n-K} \sum_{k=1}^K \sum_{i:y_i=k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$
- $\hat{\Sigma}_k = \frac{1}{n_k-K} \sum_{i:y_i=k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$
- $\pi_k = n_k/n$

To simplify calculation we use the eigenvalue decomposition of the covariance matrices

$$\hat{\Sigma}_k = \mathbf{U}_k \mathbf{D}_k \mathbf{U}_k^T$$

- \mathbf{U}_k is a $p \times p$ orthonormal matrix
- \mathbf{D}_k is a diagonal matrix of decreasing positive eigenvalues d_{kl}

The main terms in the discriminants,

$$\delta_k(x) = -\frac{1}{2} \log |\hat{\Sigma}_k| - \frac{1}{2} (x - \mu_k)^T \hat{\Sigma}_k^{-1} (x - \mu_k) + \log \pi_k$$

then turn into

$$\log |\hat{\Sigma}_k| = \sum_l \log d_{kl}$$

$$(x - \hat{\mu}_k)^T \hat{\Sigma}_k^{-1} (x - \hat{\mu}_k) = [\mathbf{U}_k^T (x - \hat{\mu}_k)]^T \mathbf{D}_k^{-1} [\mathbf{U}_k^T (x - \hat{\mu}_k)]$$

The LDA estimator

- Step 1: Normalize \mathbf{X} to spherical covariance

$$\mathbf{X}^* \leftarrow \mathbf{D}^{-1/2} \mathbf{U}^T \mathbf{X}$$

- Step 2: Classify to the closest class centroid in the transformed space, where distance is weighted by the class prior probabilities π_k

Types of Errors – a handy guide



Type I error
(false positive)



Type II error
(false negative)

Example on Multivariate LDA

Example **default** with **balance** and **student** as inputs

- training error for LDA is 2.75%
- data is highly unbalanced, we have only 3,33% positives
- the **No**-only classifier has an error of already only 3,33%

Sensitivity $\text{Sens} = TP / (TP + FN) = TP / P^*$

- fraction of correctly predicted positives

Specificity $\text{Spec} = TN / (TN + FP) = TN / N^*$

- fraction of correctly predicted negatives
- **No-only** $\text{Sens} = \frac{0}{333} = 0\%$, $\text{Spec} = \frac{9,667}{9,667} = 100\%$
- LDA $\text{Sens} = \frac{81}{333} = 24.3\%$, $\text{Spec} = \frac{9,644}{9,667} = 99.8\%$
- LDA approximates the Bayes classifier, it minimizes error on **all observations**

LDA Model Results

Prediction	True Default Status		Total
	No (−)	Yes (+)	
No (−)	9,644	252	9,896
Yes (+)	23	81	104
Total	9,667	333	10,000

Prediction	True Default Status		Total
	No (−)	Yes (+)	
No (−)	TN	FN	N
Yes (+)	FP	TP	P
Total	N*	P*	n

Type-1 error
False positive

Confusion matrix

Type-2 error
False negative

Example on Multivariate LDA

Biasing the classifier trades sensitivity for specificity

$$\log((p_k(x))/(p_l(x))) = \delta_k(x) - \delta_l(x)$$

- move the decision threshold between class **no** or **yes** away from

$$\Pr(\text{default} = \text{yes} \mid X = x) = P(Y = 1 \mid X) = 0.5$$

- we can increase sensitivity by choosing

$$\text{threshold} < 0.5$$

as this assigns more points to positive class **yes**

- for $\Pr(\text{default} = \text{yes} \mid X = x) > 0.2$

- Sens = $195/333 = 58.6\%$
- Spec = $9,432/9,667 = 97.6\%$
- Error = $373/10,000 = 3.73\%$

For a threshold of 0.5 we get
Sens = 24.3%, Spec = 99.8%, Error=2.75%

Prediction	True Default Status		
	No (−)	Yes (+)	Total
No (−)	9,644	252	9,896
Yes (+)	23	81	104
Total	9,667	333	10,000

Prediction	True Default Status		
	No (−)	Yes (+)	Total
No (−)	9,432	138	9,570
Yes (+)	235	195	430
Total	9,667	333	10,000


While for a threshold of 0.2 we have
Sens = 58.6%, Spec = 97.6%, Error=3.73%

Example on Multivariate LDA

Biassing the classifier trades sensitivity for specificity

$$\log((p_k(x))/(p_l(x))) = \delta_k(x) - \delta_l(x)$$

- move the decision threshold between class **no** or **yes** away from
 $\Pr(\text{default} = \text{yes} \mid X = x) = P(Y = 1 \mid X) = 0.5$
- we can increase sensitivity by choosing
 $\text{threshold} < 0.5$
as this assigns more points to positive class **yes**
- for $\Pr(\text{default} = \text{yes} \mid X = x) > 0.2$
 - $\text{Sens} = 195/333 = 58.6\%$
 - $\text{Spec} = 9,432/9,667 = 97.6\%$
 - $\text{Error} = 373/10,000 = 3.73\%$
- error rates change smoothly when we move the threshold



ROC Curves

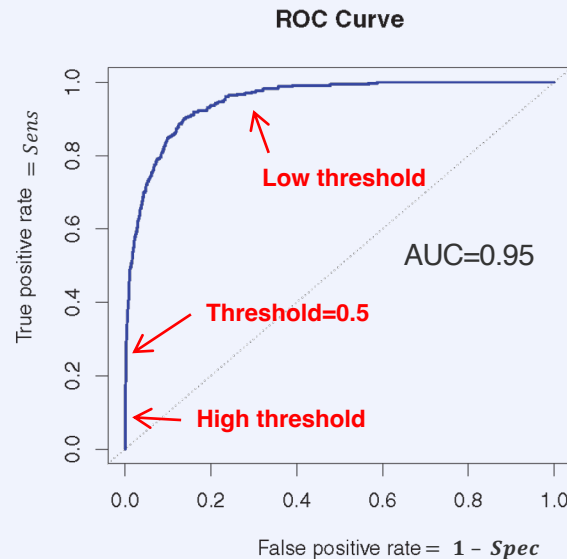
Receiver-Operating Characteristic (ROC) curves

plot *Sens* against $1 - \textit{Spec}$ for all thresholds

- Area Under the ROC-Curve (AUC) measures the quality of a classifier **independent** of the choice of that threshold
- optimally $\textit{Spec} = \textit{Sens} = 1$ for any threshold ($AUC = 1$)
- random classifier performs on the diagonal ($AUC = 0.5$)
- if the ROC curve goes below the diagonal, we can improve accuracy by inverting the classifier

ROC curves are **not influenced by imbalance** of the data

- balance only affects **locations** of a threshold along the curve



Comparing Different Classifiers

Comparison of the Classification Methods

We now know four classifiers: k -NN (L01), LDA, QDA and logistic regression

- when should we use which?

Logistic regression and LDA are surprisingly closely related

- univariate binary setting $p_2(x) = 1 - p_1(x)$
- log-odds for LDA are $\log \frac{p_1(x)}{1-p_1(x)} = c_0 + c_1x$
(difference of two linear discriminants)
- while for logistic regression $\log \frac{p_1(x)}{1-p_1(x)} = \beta_0 + \beta_1x$

Similar, but different

- β_0 and β_1 are maximum likelihood estimates
- c_0 and c_1 are estimated from sample mean and variance of Gaussian distribution
- relationship extends to multivariate data: LR and LDA often give similar results – but not always!
- LDA makes stronger assumptions (i.e., Gaussian class-conditional density)

Comparison of the Classification Methods

We now know four classifiers: k -NN, LDA, QDA and logistic regression

- when should we use which?

k -NN is nonparametric and tends to work better for strongly nonlinear settings

- it does not allow for inference, i.e. we do not get a model that we can learn from

QDA is a compromise between LDA and k -NN

Logistic regression very often works great in practice, and one can transform the features to have non-linear classification with respect to original features. Often used as baseline!


Comparing the Classification Methods

Scenario 1

- 100 random training data sets, $p = 2$ predictors, $K = 2$ classes
- 20 observations per class
- observations in different classes uncorrelated normal variables with different means and the same variance (spherical Gaussian)
- this matches the LDA assumptions of LDA

Observations

- **LDA works very well**
- logistic regression assumes a linear decision boundary, performs only slightly worse than LDA
- k -NN overfits, as does QDA




Comparing the Classification Methods

Scenario 2

- 100 random training data sets, $p = 2$ predictors, $K = 2$ classes
- like scenario 1, but predictors in each class now have a correlation of -0.5 (elliptical multivariate Gaussian)

Observations

- relative performances are similar to scenario 1 but with QDA competing as we match its assumptions.




Comparing the Classification Methods

Scenario 3

- 100 random training data sets, $p = 2$ predictors, $K = 2$ classes
- X_1 and X_2 are generated using a **t-distribution**
- more extreme points than with a Gaussian
- decision boundary is linear, but, setup violates LDA assumption

Observations

- logistic regression performs best
- QDA deteriorates because of **non-normality** of the data




Comparing the Classification Methods

Scenario 4

- 100 random training data sets, $p = 2$ predictors, $K = 2$ classes
- class 1: normal distribution with correlation **0.5** to predictors
- class 2: normal distribution with correlation **-0.5** to predictors
- assumptions of **QDA** are met (but not LDA!)

Observations

- QDA outperforms all other methods




Comparing the Classification Methods

Scenario 5

- 100 random training data sets, $p = 2$ predictors, $K = 2$ classes
- two normal distributions with uncorrelated predictors
- inputs X_1^2, X_2^2 and X_1X_2 , not X_1 and X_2
- the decision boundary is **quadratic**

Observations

- QDA performs best
- k NN (CV) follows closely
- the linear methods all perform poorly



Comparing the Classification Methods

Scenario 6

- 100 random training data sets, $p = 2$ predictors, $K = 2$ classes
- like scenario 5, but responses sampled from a complicated linear function

Observations

- even QDA cannot model data well
- k -NN-1 overfits
- k -NN (cross-validated) outperforms all parametric approaches
- smoothness must be chosen carefully

